

# Frequent Pattern Mining - Association Rules

author: IE 2064 Data Science date: April 2016

- packages used
- arules
- arulesViz
- pmml

## Introduction

type: section

## Why use Frequent Pattern Mining?

- One common purpose of data mining is to discover novel patterns in the data.
- How can we determine if elements in the data are related?
- *Association Rules* are one example of an *Unsupervised learning* method.
- You have not provided the procedure with examples of what *correct* answers look like, so the method needs to search for candidate correct answers.

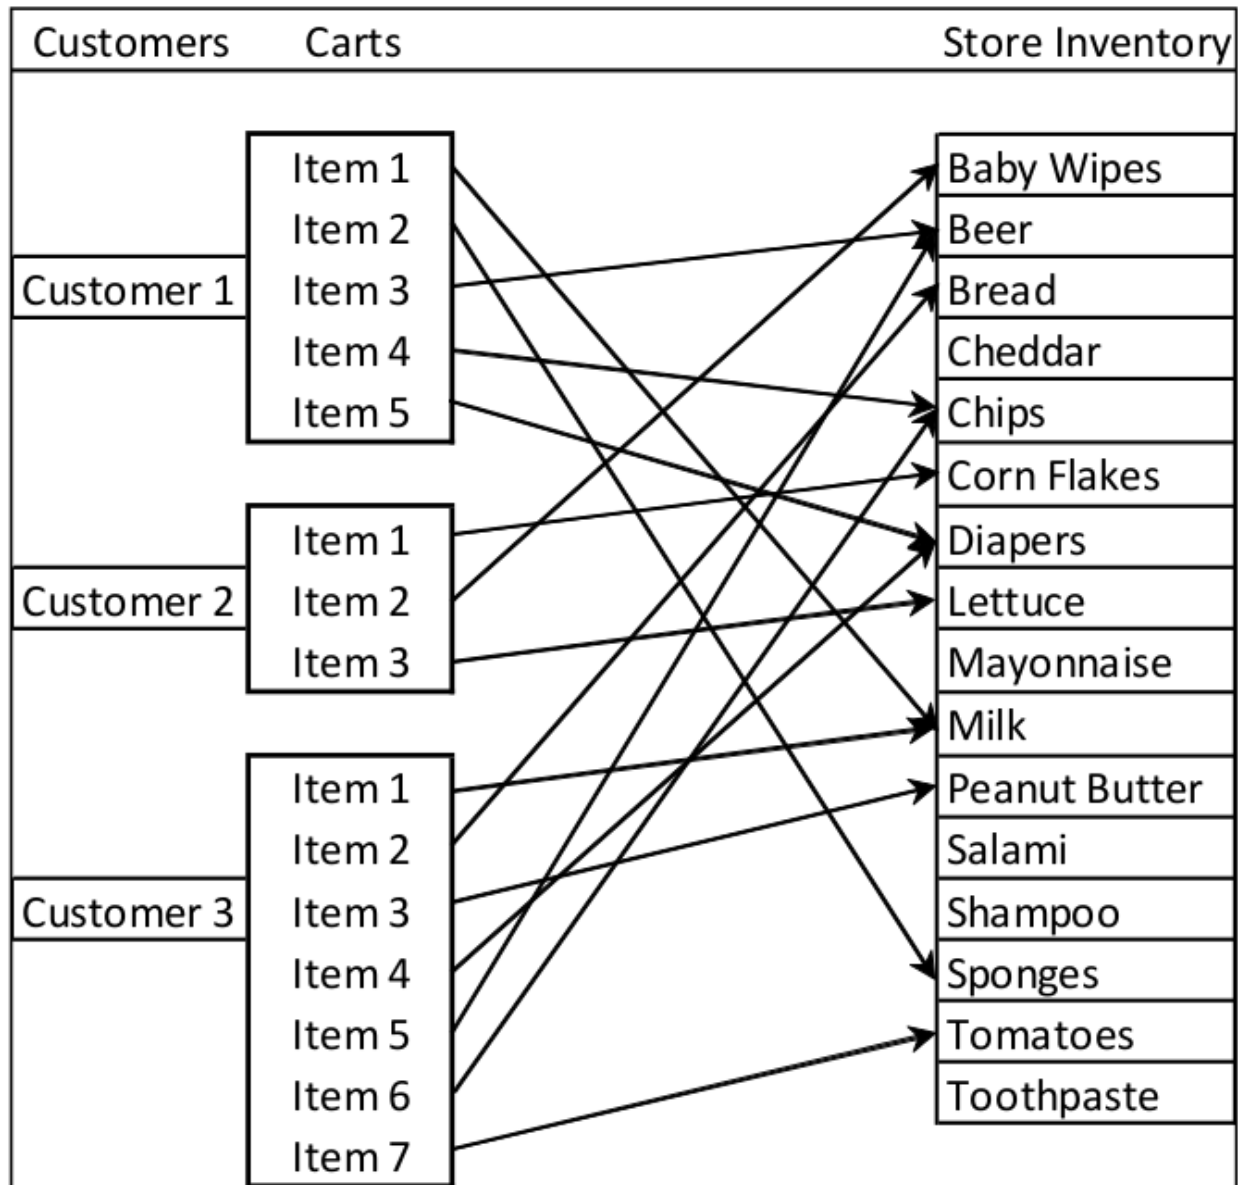
## Association Rules

type: section

## Association Rules Illustrated

left: 40% - Grocery customers and shopping baskets. - What items are commonly bought together? - If you were told an answer, what would you want to know about it?

---



## Affinity analysis

- Consider many possible propositions of combinations (rules).
- Evaluate the database of transactions to evaluate a list of rules for *support*.
- *support* - portion of cases that particular pair appears.
- *confidence* - of the cases where one member appears, the portion of the time where the second member of a pair appears.

## Groceries example

Let's look at some data.

```
library(arules)
data(Groceries)
```

## Grocery summary

```
summary(Groceries)
```

transactions as itemMatrix in sparse format with  
9835 rows (elements/itemsets/transactions) and  
169 columns (items) and a density of 0.02609146

most frequent items:

whole milk	other vegetables	rolls/buns	soda
2513	1903	1809	1715
yogurt	(Other)		
1372	34055		

element (itemset/transaction) length distribution:  
sizes

1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
2159	1643	1299	1005	855	645	545	438	350	246	182	117	78	77	55
16	17	18	19	20	21	22	23	24	26	27	28	29	32	
46	29	14	14	9	11	4	6	1	1	1	1	3	1	

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
1.000	2.000	3.000	4.409	6.000	32.000

includes extended item information - examples:

	labels	level2	level1
1	frankfurter	sausage meat	and sausage
2	sausage	sausage meat	and sausage
3	liver loaf	sausage meat	and sausage

## Data structure


- Sparse matrix
- Which items are most frequent?
- How many items in a cart?

## Itemset matrix example

		items			
		$i_1$	$i_2$	$i_3$	$i_4$
		milk	bread	butter	beer
itemsets	$X_1$	1	1	0	0
	$X_2$	0	1	0	1
	$X_3$	1	1	1	0
	$X_4$	0	0	1	0

## Some common items

```
itemFrequencyPlot(Groceries,support=0.1,cex.names=0.5)
```



## Association rules algorithms

- `apriori()`
- `eclat()`
- Parameter sets
- **parameter** changes the characteristics of the ruleset (e.g. *support*, *confidence*, *maxlen*)
- **control** influences the performance (e.g. sorting)
- **appearance** - Any restrictions
- Changing parameter values changes the results (size of subsets, number of rules generated tai)

## apriori

```
ruleset1 <-apriori(Groceries,parameter=list(support=0.005, confidence=0.5))
```

Apriori

Parameter specification:

confidence	minval	smax	arem	aval	originalSupport	support	minlen	maxlen
0.5	0.1	1	none	FALSE	TRUE	0.005	1	10
target	ext							
rules	FALSE							

Algorithmic control:

filter	tree	heap	memopt	load	sort	verbose
0.1	TRUE	TRUE	FALSE	TRUE	2	TRUE

Absolute minimum support count: 49

```
set item appearances ...[0 item(s)] done [0.00s].
set transactions ...[169 item(s), 9835 transaction(s)] done [0.00s].
sorting and recoding items ... [120 item(s)] done [0.00s].
creating transaction tree ... done [0.00s].
checking subsets of size 1 2 3 4 done [0.00s].
writing ... [120 rule(s)] done [0.00s].
creating S4 object ... done [0.00s].
```

## Reduce the number of rules

```
ruleset2 <- apriori(Groceries,parameter=list(support=0.01, confidence=0.5))
```

Apriori

Parameter specification:

confidence	minval	smax	arem	aval	originalSupport	support	minlen	maxlen
0.5	0.1	1	none	FALSE	TRUE	0.01	1	10
target	ext							
rules	FALSE							

Algorithmic control:

filter	tree	heap	memopt	load	sort	verbose
0.1	TRUE	TRUE	FALSE	TRUE	2	TRUE

Absolute minimum support count: 98

```
set item appearances ...[0 item(s)] done [0.00s].
set transactions ...[169 item(s), 9835 transaction(s)] done [0.00s].
sorting and recoding items ... [88 item(s)] done [0.00s].
creating transaction tree ... done [0.00s].
checking subsets of size 1 2 3 4 done [0.00s].
```

```
writing ... [15 rule(s)] done [0.00s].
creating S4 object ... done [0.00s].
```

## Look at the result

```
summary(ruleset2)
```

```
set of 15 rules
```

```
rule length distribution (lhs + rhs):sizes
3
15
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
3	3	3	3	3	3

```
summary of quality measures:
```

support	confidence	lift
Min. :0.01007	Min. :0.5000	Min. :1.984
1st Qu.:0.01174	1st Qu.:0.5151	1st Qu.:2.036
Median :0.01230	Median :0.5245	Median :2.203
Mean :0.01316	Mean :0.5411	Mean :2.299
3rd Qu.:0.01403	3rd Qu.:0.5718	3rd Qu.:2.432
Max. :0.02227	Max. :0.5862	Max. :3.030

```
mining info:
```

data	ntransactions	support	confidence
Groceries	9835	0.01	0.5

## lift

- How do you determine how *interesting* a rule is?
- A measure of *support* for a rule.
- Gives increased weight where the Left Hand Side or Right Hand Side occur rarely, but when they do occur, occur together.
- Larger lift is more *interesting*

## Take a closer look at the results

```
inspect(ruleset2)
```

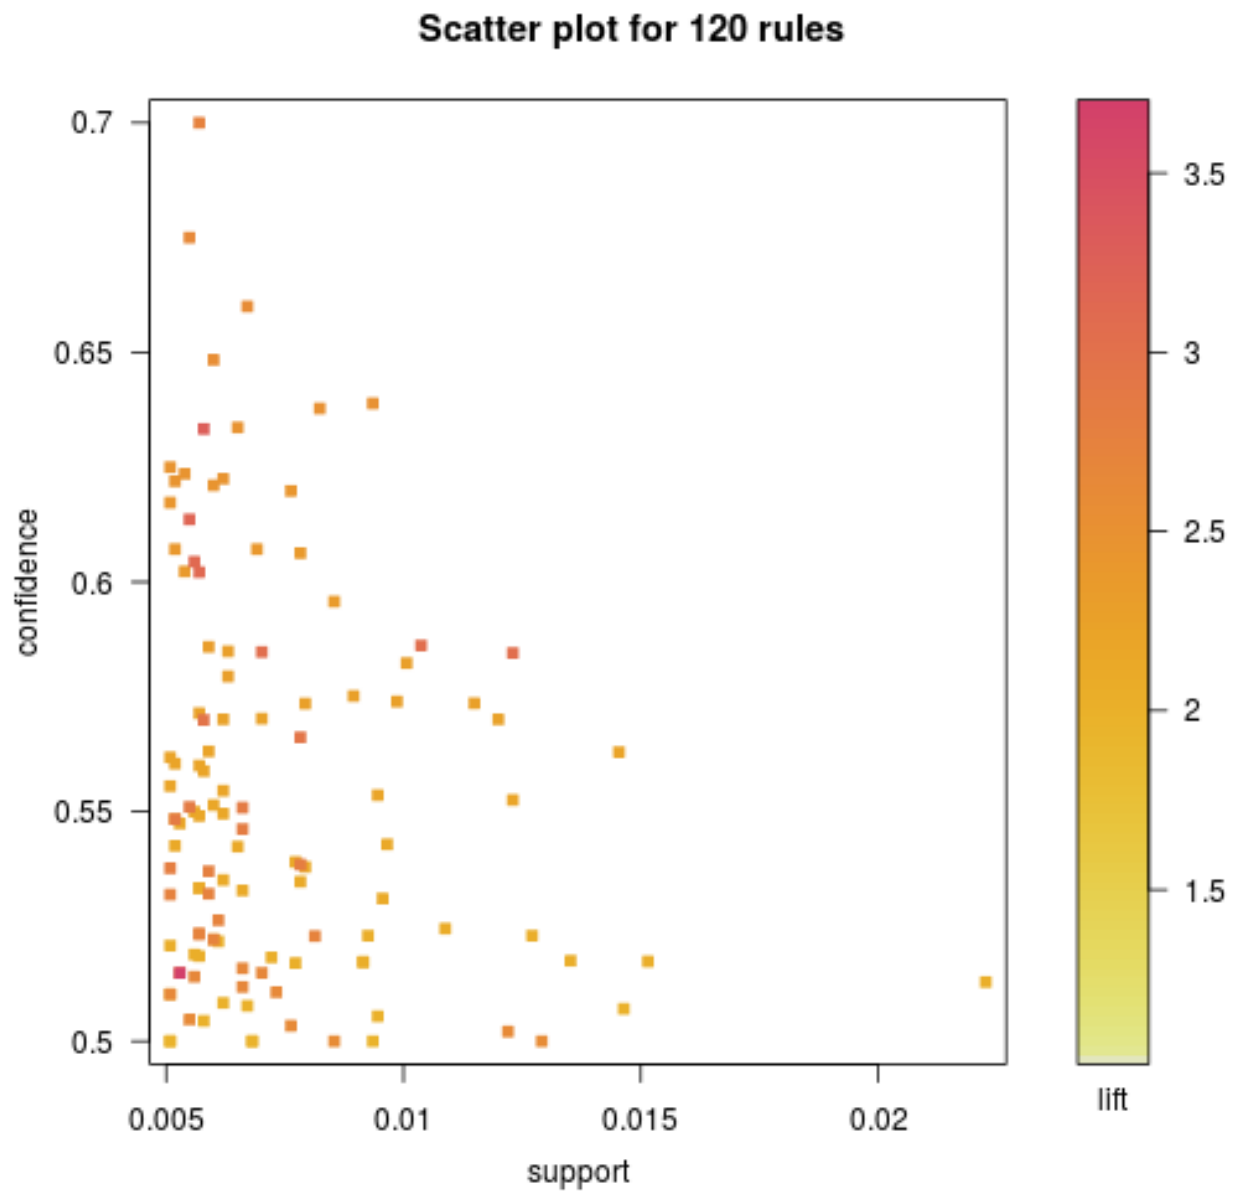
	lhs	rhs	support	confidence	lift
1	{curd, yogurt}	=> {whole milk}	0.01006609	0.5823529	2.279125
2	{other vegetables, butter}	=> {whole milk}	0.01148958	0.5736041	2.244885

3	{other vegetables, domestic eggs}	=> {whole milk}	0.01230300	0.5525114	2.162336
4	{yogurt, whipped/sour cream}	=> {whole milk}	0.01087951	0.5245098	2.052747
5	{other vegetables, whipped/sour cream}	=> {whole milk}	0.01464159	0.5070423	1.984385
6	{pip fruit, other vegetables}	=> {whole milk}	0.01352313	0.5175097	2.025351
7	{citrus fruit, root vegetables}	=> {other vegetables}	0.01037112	0.5862069	3.029608
8	{tropical fruit, root vegetables}	=> {other vegetables}	0.01230300	0.5845411	3.020999
9	{tropical fruit, root vegetables}	=> {whole milk}	0.01199797	0.5700483	2.230969
10	{tropical fruit, yogurt}	=> {whole milk}	0.01514997	0.5173611	2.024770
11	{root vegetables, yogurt}	=> {other vegetables}	0.01291307	0.5000000	2.584078
12	{root vegetables, yogurt}	=> {whole milk}	0.01453991	0.5629921	2.203354
13	{root vegetables, rolls/buns}	=> {other vegetables}	0.01220132	0.5020921	2.594890
14	{root vegetables, rolls/buns}	=> {whole milk}	0.01270971	0.5230126	2.046888
15	{other vegetables, yogurt}	=> {whole milk}	0.02226741	0.5128806	2.007235

Now for a visual inspection of results

```
library(arulesViz)
plot(ruleset1)
```





See if there are any good rules from the larger set

```
goodrules <- ruleset1[quality(ruleset1)$lift > 3.0]
inspect(goodrules)
```

	lhs	rhs	support	confidence	lift
1	{root vegetables, onions}	=> {other vegetables}	0.005693950	0.6021505	3.112008
2	{tropical fruit, curd}	=> {yogurt}	0.005287239	0.5148515	3.690645
3	{pip fruit,				

```

    whipped/sour cream} => {other vegetables} 0.005592272 0.6043956 3.123610
4 {citrus fruit,
   root vegetables}    => {other vegetables} 0.010371124 0.5862069 3.029608
5 {tropical fruit,
   root vegetables}    => {other vegetables} 0.012302999 0.5845411 3.020999
6 {pip fruit,
   root vegetables,
   whole milk}         => {other vegetables} 0.005490595 0.6136364 3.171368
7 {citrus fruit,
   root vegetables,
   whole milk}         => {other vegetables} 0.005795628 0.6333333 3.273165
8 {tropical fruit,
   root vegetables,
   whole milk}         => {other vegetables} 0.007015760 0.5847458 3.022057

```

## Data preparation example

type:section

## Epub downloads

Electronic book downloads from Vienna University of Economics

```
data(Epub)
Epub
```

```

transactions in sparse format with
15729 transactions (rows) and
936 items (columns)

```

## Get more information

```
summary(Epub)
```

```

transactions as itemMatrix in sparse format with
15729 rows (elements/itemsets/transactions) and
936 columns (items) and a density of 0.001758755

```

most frequent items:

```

doc_11d doc_813 doc_4c6 doc_955 doc_698 (Other)
   356      329      288      282      245   24393

```

```

element (itemset/transaction) length distribution:
sizes

```

```

   1     2     3     4     5     6     7     8     9    10    11    12
11615 2189  854  409  198  121   93   50   42   34   26   12
   13    14    15    16    17    18    19    20    21    22    23    24

```

10	10	6	8	6	5	8	2	2	3	2	3
25	26	27	28	30	34	36	38	41	43	52	58
4	5	1	1	1	2	1	2	1	1	1	1

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
1.000	1.000	1.000	1.646	2.000	58.000

includes extended item information - examples:

```
labels
1 doc_11d
2 doc_13d
3 doc_14c
```

includes extended transaction information - examples:

```
transactionID      TimeStamp
10792 session_4795 2003-01-01 20:59:00
10793 session_4797 2003-01-02 07:46:01
10794 session_479a 2003-01-02 10:50:38
```

## See how it changes over time

```
year <- strptime(as.POSIXlt(transactionInfo(Epub)[["TimeStamp"]]), "%Y")
table(year)
```


```
year
2003 2004 2005 2006 2007 2008
 987 1375 1611 3015 4050 4691
```

## Look at one years worth of downloads

```
epub2003 <- Epub[year=="2003"]
length(epub2003)
```

```
[1] 987
```

```
image(epub2003)
```



Let's look at only long transactions

```
transactionInfo(epub2003[size(epub2003) > 20])
```

	transactionID	TimeStamp
11092	session_56e2	2003-04-29 13:30:38
11371	session_6308	2003-08-17 18:16:12

## Let's take a closer look

```
inspect(epub2003[1:5])
```

	items	transactionID	TimeStamp
10792	{doc_154}	session_4795	2003-01-01 20:59:00
10793	{doc_3d6}	session_4797	2003-01-02 07:46:01
10794	{doc_16f}	session_479a	2003-01-02 10:50:38
10795	{doc_11d,doc_1a7,doc_f4}	session_47b7	2003-01-02 18:55:50
10796	{doc_83}	session_47bb	2003-01-02 21:27:44

## What if I want transactions per document

Coerce into a vertical layout with transaction ID list for each document.

```
epubTidLists <- as(Epub, "tidLists")  
as(epubTidLists[5], 'list')
```

```
$doc_150  
[1] "session_56e2" "session_575d" "session_7090" "session_80ef"  
[5] "session_9b5a" "session_bf41" "session_112a9" "session_11e26"  
[9] "session_123bc" "session_12938" "session_12a5e" "session_14ae7"  
[13] "session_15e17" "session_161ca" "session_177cf" "session_18649"  
[17] "session_18a83" "session_190bf" "session_19152" "session_19c27"  
[21] "session_1a264" "session_1c2e6" "session_1e935" "session_20955"  
[25] "session_23fe8"
```

## Questionnaire data example

- Source: 1994 U.S. Census
- 48842 records
- Filtered so that AAGE>16 and AGI>100
- Adults with non-zero income
- Can we determine if

```
data("AdultUCI")  
dim(AdultUCI)
```

```
[1] 48842    15
```

## Data summary

```
summary(AdultUCI)
```

age	workclass	fnlwgt
Min. :17.00	Private :33906	Min. : 12285
1st Qu.:28.00	Self-emp-not-inc: 3862	1st Qu.: 117550
Median :37.00	Local-gov : 3136	Median : 178144
Mean :38.64	State-gov : 1981	Mean : 189664
3rd Qu.:48.00	Self-emp-inc : 1695	3rd Qu.: 237642
Max. :90.00	(Other) : 1463	Max. :1490400
	NA's : 2799	

education	education-num	marital-status
HS-grad :15784	Min. : 1.00	Divorced : 6633
Some-college:10878	1st Qu.: 9.00	Married-AF-spouse : 37
Bachelors : 8025	Median :10.00	Married-civ-spouse :22379
Masters : 2657	Mean :10.08	Married-spouse-absent: 628
Assoc-voc : 2061	3rd Qu.:12.00	Never-married :16117
11th : 1812	Max. :16.00	Separated : 1530
(Other) : 7625		Widowed : 1518

occupation	relationship	race
Prof-specialty : 6172	Husband :19716	Amer-Indian-Eskimo: 470
Craft-repair : 6112	Not-in-family :12583	Asian-Pac-Islander: 1519
Exec-managerial: 6086	Other-relative: 1506	Black : 4685
Adm-clerical : 5611	Own-child : 7581	Other : 406
Sales : 5504	Unmarried : 5125	White :41762
(Other) :16548	Wife : 2331	
NA's : 2809		

sex	capital-gain	capital-loss	hours-per-week
Female:16192	Min. : 0	Min. : 0.0	Min. : 1.00
Male :32650	1st Qu.: 0	1st Qu.: 0.0	1st Qu.:40.00
	Median : 0	Median : 0.0	Median :40.00
	Mean : 1079	Mean : 87.5	Mean :40.42
	3rd Qu.: 0	3rd Qu.: 0.0	3rd Qu.:45.00
	Max. :99999	Max. :4356.0	Max. :99.00

native-country	income
United-States:43832	small:24720
Mexico : 951	large: 7841
Philippines : 295	NA's :16281
Germany : 206	
Puerto-Rico : 184	
(Other) : 2517	
NA's : 857	

## Take a closer look

```
AdultUCI[1:2,]
```

	age	workclass	fnlwgt	education	education-num	marital-status
1	39	State-gov	77516	Bachelors	13	Never-married
2	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse

	occupation	relationship	race	sex	capital-gain	capital-loss
1	Adm-clerical	Not-in-family	White	Male	2174	0
2	Exec-managerial	Husband	White	Male	0	0

```

hours-per-week native-country income
1             40 United-States small
2             13 United-States small

```

## Clean data

Remove a weighting calculation and a duplicate education factor

```

AdultUCI[["fnlwgt"]] <- NULL
AdultUCI[["education-num"]] <- NULL

```

## Map some other values to categorical variables

```

AdultUCI[["age"]] <- ordered(cut(AdultUCI[["age"]],
                                c(15,25,45,65,100)), labels = c("young", "middle-aged", "old"))
AdultUCI[["hours-per-week"]] <- ordered(cut(AdultUCI[["hours-per-week"]], c(0,25,40,60,168)), labels = c("less than 25 hrs/week", "25 to 40 hrs/week", "40 to 60 hrs/week", "more than 60 hrs/week"))
AdultUCI[["capital-gain"]] <- ordered(cut(AdultUCI[["capital-gain"]], c(-Inf,0,median(AdultUCI[["capital-gain"]]),Inf)), labels = c("less than 0", "0 to median", "more than median"))
AdultUCI[["capital-loss"]] <- ordered(cut(AdultUCI[["capital-loss"]], c(-Inf,0,median(AdultUCI[["capital-loss"]]),Inf)), labels = c("less than 0", "0 to median", "more than median"))

```

## Convert to a binary incidence matrix through coercion to transactions

```

Adult <- as(AdultUCI, "transactions")
Adult

```

```

transactions in sparse format with
48842 transactions (rows) and
115 items (columns)

```

## See what we have

```
summary(Adult)
```

```

transactions as itemMatrix in sparse format with
48842 rows (elements/itemsets/transactions) and
115 columns (items) and a density of 0.1089939

```

most frequent items:

capital-loss=none	capital-gain=None
46560	44807
native-country=United-States	race=White
43832	41762

workclass=Private	(Other)
33906	401333

element (itemset/transaction) length distribution:  
 sizes

9	10	11	12	13
19	971	2067	15623	30162

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
9.00	12.00	13.00	12.53	13.00	13.00

includes extended item information - examples:

	labels	variables	levels
1	age=Young	age	Young
2	age=Middle-aged	age	Middle-aged
3	age=Senior	age	Senior


includes extended transaction information - examples:

	transactionID
1	1
2	2
3	3

## Now plot the Item Frequency Plot

```
itemFrequencyPlot(Adult, support = 0.2, cex.names=0.8)
```





## Generate some rules

```
rules <- apriori(Adult,
  parameter =
    list(support = 0.01, confidence = 0.6))
```

Apriori

Parameter specification:

confidence	minval	smax	arem	aval	originalSupport	support	minlen	maxlen
0.6	0.1	1	none	FALSE	TRUE	0.01	1	10

```
target    ext
rules FALSE
```

Algorithmic control:

```
filter tree heap memopt load sort verbose
0.1 TRUE TRUE FALSE TRUE 2 TRUE
```

Absolute minimum support count: 488

```
set item appearances ...[0 item(s)] done [0.00s].
set transactions ...[115 item(s), 48842 transaction(s)] done [0.04s].
sorting and recoding items ... [67 item(s)] done [0.01s].
creating transaction tree ... done [0.04s].
checking subsets of size 1 2 3 4 5 6 7 8 9 10 done [1.20s].
writing ... [276443 rule(s)] done [0.04s].
creating S4 object ... done [0.22s].
```

## Summarize the rules

```
summary(rules)
```

set of 276443 rules

rule length distribution (lhs + rhs):sizes

1	2	3	4	5	6	7	8	9	10
6	432	4981	22127	52669	75104	67198	38094	13244	2588

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
1.000	5.000	6.000	6.289	7.000	10.000

summary of quality measures:

support	confidence	lift
Min. :0.01001	Min. :0.6000	Min. : 0.7171
1st Qu.:0.01253	1st Qu.:0.7691	1st Qu.: 1.0100
Median :0.01701	Median :0.9051	Median : 1.0554
Mean :0.02679	Mean :0.8600	Mean : 1.3109
3rd Qu.:0.02741	3rd Qu.:0.9542	3rd Qu.: 1.2980
Max. :0.95328	Max. :1.0000	Max. :20.6826

mining info:

data	ntransactions	support	confidence
Adult	48842	0.01	0.6

## Break data into subset, and limit the number of rules

- Create rules for both ‘income-small’ and ‘income-large’
- Limit the number of rules by specifying a minimum lift.

```
rulesIncomeSmall <- subset(rules, subset =
  rhs %in% "income=small" & lift > 1.2)
rulesIncomeLarge <- subset(rules, subset =
  rhs %in% "income=large" & lift > 1.2)
```

## Inspect the best rules

```
inspect(head(sort(rulesIncomeSmall, by = "confidence"),
  n = 3))
```

	lhs	rhs	support	confidence	lift
1	{workclass=Private, marital-status=Never-married, relationship=Own-child, sex=Male, hours-per-week=Part-time, native-country=United-States}	=> {income=small}	0.01074895	0.7104195	1.403653
2	{workclass=Private, marital-status=Never-married, relationship=Own-child, sex=Male, hours-per-week=Part-time}	=> {income=small}	0.01144507	0.7102922	1.403402
3	{workclass=Private, marital-status=Never-married, relationship=Own-child, sex=Male, capital-gain=None, hours-per-week=Part-time, native-country=United-States}	=> {income=small}	0.01046231	0.7097222	1.402276

## Inspect when income large

```
inspect(head(sort(rulesIncomeLarge, by = "confidence"),
  n = 3))
```

	lhs	rhs	support	confidence	lift
1	{marital-status=Married-civ-spouse, capital-gain=High, native-country=United-States}	=> {income=large}	0.01562180	0.6849192	4.266398
2	{marital-status=Married-civ-spouse, capital-gain=High, capital-loss=none, native-country=United-States}	=> {income=large}	0.01562180	0.6849192	4.266398
3	{relationship=Husband, race=White, capital-gain=High, native-country=United-States}	=> {income=large}	0.01302158	0.6846071	4.264454

## Save the rules

Save these rules using PMML for use in other systems.

```
library(pmml)
write.PMML(rulesIncomeSmall, file = "incomerulessmall.xml")
```